Intro to Bayesian Computing

Krzysztof Latuszynski
(University of Warwick, UK)

OxWaSP - module 1

Prior-Posterior

- let $\theta \in \Theta$ be a parameter of a statistical model, say $M(\theta)$.
  E.g. $\Theta \in \mathbb{R}^d$, $\Theta \in \mathbb{N}^d$, $\Theta \in \{0,1\}^d$
- In Bayesian Statistics one assumes θ is random, i.e. there exists a prior probability distribution $p(\theta)$ on Θ s.t. in absence of additional information $\theta \sim p(\theta)$.
- $y_1, \ldots, y_n \in \mathbb{Y}^n$- data
- $l(\theta|y_1, \ldots, y_n)$ - the likelihood function for the model $M(\theta)$
- Example: Consider a diffusion model $M(\theta)$ where $\theta = (\mu, \sigma)$

$$dX_t = \mu dt + \sigma dB_t$$

  observed at discrete time points $(t_0, t_1, \ldots, t_N)$ as $(x_{t_0}, x_{t_1}, \ldots, x_{t_N})$
- The likelihood function is

$$l(\theta|x_{t_0}, x_{t_1}, \ldots, x_{t_N}) = \prod_{i=1}^{N} l(\theta|x_{t_i}, x_{t_{i-1}}) = \prod_{i=1}^{N} \phi_{N(\mu(t_i - t_{i-1}), \sigma^2(t_i - t_{i-1}))}(x_{t_i} - x_{t_{i-1}}).$$

Prior-Posterior

- let $\theta \in \Theta$ be a parameter of a statistical model, say $M(\theta)$.
  E.g. $\Theta \in \mathbb{R}^d, \quad \Theta \in \mathbb{N}^d, \quad \Theta \in \{0,1\}^d$
- In Bayesian Statistics one assumes θ is random, i.e. there exists a prior probability distribution $p(\theta)$ on Θ s.t. in absence of additional information $\theta \sim p(\theta)$.
- $y_1, \ldots, y_n \in \mathbb{Y}^n$- data
- $l(\theta|y_1, \ldots, y_n)$ - the likelihood function for the model $M(\theta)$
- Example: Consider a diffusion model $M(\theta)$ where $\theta = (\mu, \sigma)$

  $$dX_t = \mu dt + \sigma dB_t$$

  observed at discrete time points $(t_0, t_1, \ldots, t_N)$ as $(x_{t_0}, x_{t_1}, \ldots, x_{t_N})$
- The likelihood function is

$$l(\theta|x_{t_0}, x_{t_1}, \ldots, x_{t_N}) = \prod_{i=1}^{N} l(\theta|x_{t_i}, x_{t_{i-1}}) = \prod_{i=1}^{N} \phi_{N(\mu(t_i - t_{i-1}), \sigma^2(t_i - t_{i-1}))}(x_{t_i} - x_{t_{i-1}}).$$

Prior-Posterior

- let $\theta \in \Theta$ be a parameter of a statistical model, say $M(\theta)$.
  E.g. $\Theta \in \mathbb{R}^d$, $\Theta \in \mathbb{N}^d$, $\Theta \in \{0,1\}^d$
- In Bayesian Statistics one assumes θ is random, i.e. there exists a prior probability distribution $p(\theta)$ on Θ s.t. in absence of additional information $\theta \sim p(\theta)$.
- $y_1, \ldots, y_n \in \mathbb{Y}^n$- data
- $l(\theta|y_1, \ldots, y_n)$ - the likelihood function for the model $M(\theta)$
- Example: Consider a diffusion model $M(\theta)$ where $\theta = (\mu, \sigma)$

$$dX_t = \mu dt + \sigma dB_t$$

  observed at discrete time points $(t_0, t_1, \ldots, t_N)$ as $(x_{t_0}, x_{t_1}, \ldots, x_{t_N})$
- The likelihood function is

$$l(\theta|x_{t_0}, x_{t_1}, \ldots, x_{t_N}) = \prod_{i=1}^{N} l(\theta|x_{t_i}, x_{t_{i-1}}) = \prod_{i=1}^{N} \phi_{N(\mu(t_i-t_{i-1}),\sigma^2(t_i-t_{i-1}))}(x_{t_i} - x_{t_{i-1}}).$$

Prior-Posterior

- let $\theta \in \Theta$ be a parameter of a statistical model, say $M(\theta)$.
  E.g. $\Theta \in \mathbb{R}^d$, $\Theta \in \mathbb{N}^d$, $\Theta \in \{0,1\}^d$
- In Bayesian Statistics one assumes θ is random, i.e. there exists a prior probability distribution $p(\theta)$ on Θ s.t. in absence of additional information $\theta \sim p(\theta)$.
- $y_1, \ldots, y_n \in \mathbb{Y}^n$- data
- $l(\theta|y_1, \ldots, y_n)$ - the likelihood function for the model $M(\theta)$
- Example: Consider a diffusion model $M(\theta)$ where $\theta = (\mu, \sigma)$

$$dX_t = \mu dt + \sigma dB_t$$

  observed at discrete time points $(t_0, t_1, \ldots, t_N)$ as $(x_{t_0}, x_{t_1}, \ldots, x_{t_N})$
- The likelihood function is

$$l(\theta|x_{t_0}, x_{t_1}, \ldots, x_{t_N}) = \prod_{i=1}^{N} l(\theta|x_{t_i}, x_{t_{i-1}}) = \prod_{i=1}^{N} \phi_{N(\mu(t_i - t_{i-1}), \sigma^2(t_i - t_{i-1}))}(x_{t_i} - x_{t_{i-1}}).$$

Prior-Posterior

- let $\theta \in \Theta$ be a parameter of a statistical model, say $M(\theta)$.
  E.g. $\Theta \in \mathbb{R}^d$, $\Theta \in \mathbb{N}^d$, $\Theta \in \{0,1\}^d$
- In Bayesian Statistics one assumes θ is random, i.e. there exists a prior probability distribution $p(\theta)$ on Θ s.t. in absence of additional information $\theta \sim p(\theta)$.
- $y_1, \ldots, y_n \in \mathbb{Y}^n$- data
- $l(\theta|y_1, \ldots, y_n)$ - the likelihood function for the model $M(\theta)$
- Example: Consider a diffusion model $M(\theta)$ where $\theta = (\mu, \sigma)$

$$dX_t = \mu dt + \sigma dB_t$$

  observed at discrete time points $(t_0, t_1, \ldots, t_N)$ as $(x_{t_0}, x_{t_1}, \ldots, x_{t_N})$
- The likelihood function is

$$l(\theta|x_{t_0}, x_{t_1}, \ldots, x_{t_N}) = \prod_{i=1}^{N} l(\theta|x_{t_i}, x_{t_{i-1}}) = \prod_{i=1}^{N} \phi_{N(\mu(t_i - t_{i-1}), \sigma^2(t_i - t_{i-1}))}(x_{t_i} - x_{t_{i-1}}).$$

Prior-Posterior

- let $\theta \in \Theta$ be a parameter of a statistical model, say $M(\theta)$.
  E.g. $\Theta \in \mathbb{R}^d, \quad \Theta \in \mathbb{N}^d, \quad \Theta \in \{0,1\}^d$
- In Bayesian Statistics one assumes θ is random, i.e. there exists a prior probability distribution $p(\theta)$ on Θ s.t. in absence of additional information $\theta \sim p(\theta)$.
- $y_1, \ldots, y_n \in \mathbb{Y}^n$- data
- $l(\theta|y_1, \ldots, y_n)$ - the likelihood function for the model $M(\theta)$
- Example: Consider a diffusion model $M(\theta)$ where $\theta = (\mu, \sigma)$

$$dX_t = \mu dt + \sigma dB_t$$

  observed at discrete time points $(t_0, t_1, \ldots, t_N)$ as $(x_{t_0}, x_{t_1}, \ldots, x_{t_N})$
- The likelihood function is

$$l(\theta|x_{t_0}, x_{t_1}, \ldots, x_{t_N}) = \prod_{i=1}^{N} l(\theta|x_{t_i}, x_{t_{i-1}}) = \prod_{i=1}^{N} \phi_{N(\mu(t_i - t_{i-1}), \sigma^2(t_i - t_{i-1}))}(x_{t_i} - x_{t_{i-1}}).$$

Posterior and uncertainty quantification

- The posterior distribution is then

$$\pi(\theta) = \pi(\theta|y_1, \ldots, y_n) = \frac{p(\theta)l(\theta|y_1, \ldots, y_n)}{\int_\Theta p(\theta)l(\theta|y_1, \ldots, y_n)d\theta}.$$

- This posterior summarises uncertainty about the parameter $\theta \in \Theta$ and is used for all inferential questions like credible sets, decision making, prediction, model choice, etc.
- In the diffusion example predicting the value of the diffusion at time $t > t_N$ would amount to repeating the following steps:
  1. sample $\theta = (\mu, \sigma) \sim \pi(\theta)$
  2. sample $X_t \sim N(x_{t_N} + \mu(t - t_N), \sigma^2(t - t_N))$

Posterior and uncertainty quantification

- The posterior distribution is then

$$\pi(\theta) = \pi(\theta|y_1, \dots, y_n) = \frac{p(\theta)l(\theta|y_1, \dots, y_n)}{\int_\Theta p(\theta)l(\theta|y_1, \dots, y_n)d\theta}.$$

- This posterior summarises uncertainty about the parameter $\theta \in \Theta$ and is used for all inferential questions like credible sets, decision making, prediction, model choice, etc.
- In the diffusion example predicting the value of the diffusion at time $t > t_N$ would amount to repeating the following steps:
  1. sample $\theta = (\mu, \sigma) \sim \pi(\theta)$
  2. sample $X_t \sim N(x_{t_N} + \mu(t - t_N), \sigma^2(t - t_N))$

Posterior and uncertainty quantification

- The posterior distribution is then

$$\pi(\theta) = \pi(\theta|y_1, \ldots, y_n) = \frac{p(\theta)l(\theta|y_1, \ldots, y_n)}{\int_\Theta p(\theta)l(\theta|y_1, \ldots, y_n)d\theta}.$$

- This posterior summarises uncertainty about the parameter $\theta \in \Theta$ and is used for all inferential questions like credible sets, decision making, prediction, model choice, etc.
- In the diffusion example predicting the value of the diffusion at time $t > t_N$ would amount to repeating the following steps:
  1. sample $\theta = (\mu, \sigma) \sim \pi(\theta)$
  2. sample $X_t \sim N(x_{t_N} + \mu(t - t_N), \sigma^2(t - t_N))$

the MAP estimator

- One of the classical estimation tasks is to compute the Maximum a Posteriori Estimator (MAP), say θ_{MAP}.
- $$\theta_{MAP} := \text{argmax}_\theta \pi(\theta) = \text{argmax}_\theta \Big\{ p(\theta) l(\theta | y_1, \ldots, y_n) \Big\}$$
- Computing θ_{MAP} may be nontrivial, especially if $\pi(\theta)$ is multimodal.
- There are specialised algorithms for doing this.
- Some non-bayesian statistical inference approaches can be rewritten as bayesian MAP estimators (for example the LASSO).

the MAP estimator

- One of the classical estimation tasks is to compute the Maximum a Posteriori Estimator (MAP), say θ_{MAP}.
- $$\theta_{MAP} := \mathrm{argmax}_\theta \pi(\theta) = \mathrm{argmax}_\theta \Big\{ p(\theta) l(\theta | y_1, \ldots, y_n) \Big\}$$
- Computing θ_{MAP} may be nontrivial, especially if $\pi(\theta)$ is multimodal.
- There are specialised algorithms for doing this.
- Some non-bayesian statistical inference approaches can be rewritten as bayesian MAP estimators (for example the LASSO).

the MAP estimator

- One of the classical estimation tasks is to compute the Maximum a Posteriori Estimator (MAP), say θ_{MAP}.
- $$\theta_{MAP} := \mathrm{argmax}_\theta \pi(\theta) = \mathrm{argmax}_\theta \Big\{ p(\theta) l(\theta | y_1, \ldots, y_n) \Big\}$$
- Computing θ_{MAP} may be nontrivial, especially if $\pi(\theta)$ is multimodal.
- There are specialised algorithms for doing this.
- Some non-bayesian statistical inference approaches can be rewritten as bayesian MAP estimators (for example the LASSO).

the MAP estimator

- One of the classical estimation tasks is to compute the Maximum a Posteriori Estimator (MAP), say θ_{MAP}.
- $$\theta_{MAP} := \mathrm{argmax}_\theta \pi(\theta) = \mathrm{argmax}_\theta \Big\{ p(\theta) l(\theta | y_1, \ldots, y_n) \Big\}$$
- Computing θ_{MAP} may be nontrivial, especially if $\pi(\theta)$ is multimodal.
- There are specialised algorithms for doing this.
- Some non-bayesian statistical inference approaches can be rewritten as bayesian MAP estimators (for example the LASSO).

the MAP estimator

- One of the classical estimation tasks is to compute the Maximum a Posteriori Estimator (MAP), say θ_{MAP}.
- $$\theta_{MAP} := \text{argmax}_\theta \pi(\theta) = \text{argmax}_\theta \Big\{ p(\theta) l(\theta | y_1, \ldots, y_n) \Big\}$$
- Computing θ_{MAP} may be nontrivial, especially if $\pi(\theta)$ is multimodal.
- There are specialised algorithms for doing this.
- Some non-bayesian statistical inference approaches can be rewritten as bayesian MAP estimators (for example the LASSO).

the Bayesian estimator

- Bayesian estimator is an estimator that minimizes the posterior expected value of a loss function.
- The loss function

$$L(\cdot,\cdot) : \Theta \times \Theta \to \mathbb{R}$$

- After seeing data $(y_1, \ldots, y_n)$ we choose an estimator $\hat{\theta}(y_1, \ldots, y_n)$
- Its expected loss is

$$\begin{aligned} \mathbb{E}L(\theta, \hat{\theta}(y_1, \ldots, y_n)) &= \int_{\mathbb{Y}^n \times \Theta} L(\theta, \hat{\theta}(y_1, \ldots, y_n)) m(y_1, \ldots, y_n | \theta) p(\theta) \\ &= \int_{\mathbb{Y}^n \times \Theta} L(\theta, \hat{\theta}(y_1, \ldots, y_n)) \pi(\theta) p(dy) \end{aligned}$$

- $\hat{\theta}(y_1, \ldots, y_n)$ is a Bayesian estimator if it minimizes the above expected loss.

the Bayesian estimator

- Bayesian estimator is an estimator that minimizes the posterior expected value of a loss function.
- The loss function

$$L(\cdot,\cdot) : \Theta \times \Theta \to \mathbb{R}$$

- After seeing data $(y_1, \ldots, y_n)$ we choose an estimator $\hat{\theta}(y_1, \ldots, y_n)$
- Its expected loss is

$$\begin{aligned}\mathbb{E}L(\theta, \hat{\theta}(y_1, \ldots, y_n)) &= \int_{\mathbb{Y}^n \times \Theta} L(\theta, \hat{\theta}(y_1, \ldots, y_n)) m(y_1, \ldots, y_n | \theta) p(\theta) \\ &= \int_{\mathbb{Y}^n \times \Theta} L(\theta, \hat{\theta}(y_1, \ldots, y_n)) \pi(\theta) p(dy)\end{aligned}$$

- $\hat{\theta}(y_1, \ldots, y_n)$ is a Bayesian estimator if it minimizes the above expected loss.

the Bayesian estimator

- Bayesian estimator is an estimator that minimizes the posterior expected value of a loss function.
- The loss function
$$L(\cdot,\cdot) : \Theta \times \Theta \to \mathbb{R}$$
- After seeing data $(y_1, \ldots, y_n)$ we choose an estimator $\hat{\theta}(y_1, \ldots, y_n)$
- Its expected loss is
$$\begin{aligned}\mathbb{E}L(\theta, \hat{\theta}(y_1, \ldots, y_n)) &= \int_{\mathbb{Y}^n \times \Theta} L(\theta, \hat{\theta}(y_1, \ldots, y_n)) m(y_1, \ldots, y_n|\theta) p(\theta) \\ &= \int_{\mathbb{Y}^n \times \Theta} L(\theta, \hat{\theta}(y_1, \ldots, y_n)) \pi(\theta) p(dy)\end{aligned}$$
- $\hat{\theta}(y_1, \ldots, y_n)$ is a Bayesian estimator if it minimizes the above expected loss.

the Bayesian estimator

- Bayesian estimator is an estimator that minimizes the posterior expected value of a loss function.
- The loss function

$$L(\cdot,\cdot) : \Theta \times \Theta \to \mathbb{R}$$

- After seeing data $(y_1,\ldots,y_n)$ we choose an estimator $\hat{\theta}(y_1,\ldots,y_n)$
- Its expected loss is

$$\begin{aligned}\mathbb{E}L(\theta,\hat{\theta}(y_1,\ldots,y_n)) &= \int_{\mathbb{Y}^n\times\Theta} L(\theta,\hat{\theta}(y_1,\ldots,y_n))m(y_1,\ldots,y_n|\theta)p(\theta) \\ &= \int_{\mathbb{Y}^n\times\Theta} L(\theta,\hat{\theta}(y_1,\ldots,y_n))\pi(\theta)p(dy)\end{aligned}$$

- $\hat{\theta}(y_1,\ldots,y_n)$ is a Bayesian estimator if it minimizes the above expected loss.

the Bayesian estimator

- Bayesian estimator is an estimator that minimizes the posterior expected value of a loss function.
- The loss function
$$L(\cdot,\cdot) : \Theta \times \Theta \to \mathbb{R}$$
- After seeing data $(y_1, \ldots, y_n)$ we choose an estimator $\hat{\theta}(y_1, \ldots, y_n)$
- Its expected loss is

$$\begin{aligned}\mathbb{E}L(\theta, \hat{\theta}(y_1, \ldots, y_n)) &= \int_{\mathbb{Y}^n \times \Theta} L(\theta, \hat{\theta}(y_1, \ldots, y_n)) m(y_1, \ldots, y_n|\theta) p(\theta) \\ &= \int_{\mathbb{Y}^n \times \Theta} L(\theta, \hat{\theta}(y_1, \ldots, y_n)) \pi(\theta) p(dy)\end{aligned}$$

- $\hat{\theta}(y_1, \ldots, y_n)$ is a Bayesian estimator if it minimizes the above expected loss.

the Bayesian estimator and computing integrals

- We consider only the most common choice of quadratic loss function
$$L(\theta_1, \theta_2) = (\theta_1 - \theta_2)^2$$
- in which case
$$\hat{\theta}(y_1, \ldots, y_n) = \mathbb{E}_\pi \theta$$
so it is the posterior mean.
- So computing the Bayesian estimator is computing the integral wrt the posterior
$$\int_\Theta \theta\pi(\theta)$$
- Similarly answering other inferential questions like credible sets, posterior variance etc involve computing integrals of the form
$$\int_\Theta f(\theta)\pi(\theta).$$

the Bayesian estimator and computing integrals

- We consider only the most common choice of quadratic loss function
$$L(\theta_1, \theta_2) = (\theta_1 - \theta_2)^2$$
- in which case
$$\hat{\theta}(y_1, \ldots, y_n) = \mathbb{E}_\pi \theta$$
so it is the posterior mean.
- So computing the Bayesian estimator is computing the integral wrt the posterior
$$\int_\Theta \theta \pi(\theta)$$
- Similarly answering other inferential questions like credible sets, posterior variance etc involve computing integrals of the form
$$\int_\Theta f(\theta)\pi(\theta).$$

the Bayesian estimator and computing integrals

- We consider only the most common choice of quadratic loss function
$$L(\theta_1, \theta_2) = (\theta_1 - \theta_2)^2$$
- in which case
$$\hat{\theta}(y_1, \dots, y_n) = \mathbb{E}_\pi \theta$$
so it is the posterior mean.
- So computing the Bayesian estimator is computing the integral wrt the posterior
$$\int_\Theta \theta\pi(\theta)$$
- Similarly answering other inferential questions like credible sets, posterior variance etc involve computing integrals of the form
$$\int_\Theta f(\theta)\pi(\theta).$$

the Bayesian estimator and computing integrals

- We consider only the most common choice of quadratic loss function
$$L(\theta_1, \theta_2) = (\theta_1 - \theta_2)^2$$
- in which case
$$\hat{\theta}(y_1, \dots, y_n) = \mathbb{E}_\pi \theta$$
so it is the posterior mean.
- So computing the Bayesian estimator is computing the integral wrt the posterior
$$\int_\Theta \theta \pi(\theta)$$
- Similarly answering other inferential questions like credible sets, posterior variance etc involve computing integrals of the form
$$\int_\Theta f(\theta)\pi(\theta).$$

The Monte Carlo Method

- $$I(f) = \int_{\Theta} f(\theta)\pi(\theta).$$
- Standard Monte Carlo amounts to
  1. sample $\theta_i \sim \pi \qquad \text{for} \quad i = 1, \ldots, k$
  2. compute $\hat{I}_k(f) = \frac{1}{k}\sum_i f(\theta_i)$
- Standard LLN and CLT apply.
- In particular the CLT variance is $Var_\pi f$
-
- However sampling from π is typically not easy.

The Monte Carlo Method

- $$I(f) = \int_{\Theta} f(\theta)\pi(\theta).$$
- Standard Monte Carlo amounts to
  1. sample $\theta_i \sim \pi \qquad \text{for} \quad i = 1, \ldots, k$
  2. compute $\hat{I}_k(f) = \frac{1}{k} \sum_i f(\theta_i)$
- Standard LLN and CLT apply.
- In particular the CLT variance is $Var_{\pi} f$
-
- However sampling from π is typically not easy.

The Monte Carlo Method

- $$I(f) = \int_{\Theta} f(\theta)\pi(\theta).$$
- Standard Monte Carlo amounts to
  1. sample $\theta_i \sim \pi \qquad \text{for} \quad i = 1, \ldots, k$
  2. compute $\hat{I}_k(f) = \frac{1}{k}\sum_i f(\theta_i)$
- Standard LLN and CLT apply.
- In particular the CLT variance is $Var_{\pi} f$
-
- However sampling from π is typically not easy.

The Monte Carlo Method

- $$I(f) = \int_\Theta f(\theta)\pi(\theta).$$
- Standard Monte Carlo amounts to
  1. sample $\theta_i \sim \pi \qquad \text{for} \quad i = 1, \ldots, k$
  2. compute $\hat{I}_k(f) = \frac{1}{k}\sum_i f(\theta_i)$
- Standard LLN and CLT apply.
- In particular the CLT variance is $Var_\pi f$
-
- However sampling from π is typically not easy.

The Monte Carlo Method

- $$I(f) = \int_{\Theta} f(\theta)\pi(\theta).$$
- Standard Monte Carlo amounts to
  1. sample $\theta_i \sim \pi \qquad \text{for} \quad i = 1, \ldots, k$
  2. compute $\hat{I}_k(f) = \frac{1}{k}\sum_i f(\theta_i)$
- Standard LLN and CLT apply.
- In particular the CLT variance is $Var_{\pi} f$
-
- However sampling from π is typically not easy.

The Monte Carlo Method

- $$I(f) = \int_{\Theta} f(\theta)\pi(\theta).$$
- Standard Monte Carlo amounts to
  1. sample $\theta_i \sim \pi \qquad \text{for} \quad i = 1, \ldots, k$
  2. compute $\hat{I}_k(f) = \frac{1}{k}\sum_i f(\theta_i)$
- Standard LLN and CLT apply.
- In particular the CLT variance is $Var_{\pi}f$
-
- However sampling from π is typically not easy.

for toy distributions only

- Let F be the cdf of π and define its left continuous inverse version
- $$F^{-} := \inf\{x : F(x) \geq u\} \qquad \text{for} \quad 0 < u < 1.$$
- If $U \sim U(0,1)$ then
- $F^{-}(U) \sim \pi$
- Verify the above as an exercise.

for toy distributions only

- Let F be the cdf of π and define its left continuous inverse version
- $$F^{-} := \inf\{x : F(x) \geq u\} \qquad \text{for} \quad 0 < u < 1.$$
- If $U \sim U(0,1)$ then
- $F^{-}(U) \sim \pi$
- Verify the above as an exercise.

for toy distributions only

- Let F be the cdf of π and define its left continuous inverse version
- $$F^{-} := \inf\{x : F(x) \geq u\} \qquad \text{for} \quad 0 < u < 1.$$
- If $U \sim U(0,1)$ then
- $F^{-}(U) \sim \pi$
- Verify the above as an exercise.

for toy distributions only

- Let F be the cdf of π and define its left continuous inverse version
- $$F^{-} := \inf\{x : F(x) \geq u\} \qquad \text{for} \quad 0 < u < 1.$$
- If $U \sim U(0,1)$ then
- $F^{-}(U) \sim \pi$
- Verify the above as an exercise.

for toy distributions only

- Let F be the cdf of π and define its left continuous inverse version
- $$F^{-} := \inf\{x : F(x) \geq u\} \qquad \text{for} \quad 0 < u < 1.$$
- If $U \sim U(0,1)$ then
- $F^{-}(U) \sim \pi$
- Verify the above as an exercise.

Rejection sampling

- Sample candidate Y from density $g(\theta)$ such that

$$\pi(\theta) \leq Cg(\theta) \qquad \text{for some} \quad C < \infty$$

- accept candidate Y as θ with probability

$$\frac{\pi(Y)}{Cg(Y)}$$

  otherwise start from the beginning.
- The accepted outcome is distributed as π
- The average number of trials until acceptance is C.
- Verify the above as an exercise.

Rejection sampling

- Sample candidate Y from density $g(\theta)$ such that

$$\pi(\theta) \leq Cg(\theta) \qquad \text{for some} \quad C < \infty$$

- accept candidate Y as θ with probability

$$\frac{\pi(Y)}{Cg(Y)}$$

otherwise start from the beginning.
- The accepted outcome is distributed as π
- The average number of trials until acceptance is C.
- Verify the above as an exercise.

Rejection sampling

- Sample candidate Y from density $g(\theta)$ such that

$$\pi(\theta) \leq Cg(\theta) \qquad \text{for some} \quad C < \infty$$

- accept candidate Y as θ with probability

$$\frac{\pi(Y)}{Cg(Y)}$$

  otherwise start from the beginning.
- The accepted outcome is distributed as π
- The average number of trials until acceptance is C.
- Verify the above as an exercise.

Rejection sampling

- Sample candidate Y from density $g(\theta)$ such that

$$\pi(\theta) \le Cg(\theta) \qquad \text{for some} \quad C < \infty$$

- accept candidate Y as θ with probability

$$\frac{\pi(Y)}{Cg(Y)}$$

  otherwise start from the beginning.
- The accepted outcome is distributed as π
- The average number of trials until acceptance is C.
- Verify the above as an exercise.

Rejection sampling

- Sample candidate Y from density $g(\theta)$ such that

$$\pi(\theta) \leq Cg(\theta) \qquad \text{for some} \quad C < \infty$$

- accept candidate Y as θ with probability

$$\frac{\pi(Y)}{Cg(Y)}$$

  otherwise start from the beginning.
- The accepted outcome is distributed as π
- The average number of trials until acceptance is C.
- Verify the above as an exercise.

Importance sampling

- Let g be a density such that $\pi(\theta) > 0 \implies g(\theta) > 0$
- Then we can write

$$I = \mathbb{E}_\pi f = \int_\Theta f(\theta)\pi(\theta)d\theta = \int_\Theta f(\theta)\frac{\pi(\theta)}{g(\theta)}g(\theta)d\theta$$
$$= \int_\Theta f(\theta)W(\theta)g(\theta)d\theta = \mathbb{E}_g fW.$$

- Hence the importance sampling Algorithm:
- 1. Sample θ_i $i = 1, \ldots, k$ iid from g
  2. Estimate the integral by the unbiased, consistent estimator:

$$\hat{I}_k = \frac{1}{k}\sum_i f(\theta_i)W(\theta_i).$$

- Note that compared to iid Monte Carlo the variance of the estimators changes (typically increases) to $\mathrm{Var}_g(fW)$.

Importance sampling

- Let g be a density such that $\pi(\theta) > 0 \implies g(\theta) > 0$
- Then we can write

$$\begin{aligned} I = \mathbb{E}_\pi f &= \int_\Theta f(\theta)\pi(\theta)d\theta = \int_\Theta f(\theta)\frac{\pi(\theta)}{g(\theta)}g(\theta)d\theta \\ &= \int_\Theta f(\theta)W(\theta)g(\theta)d\theta = \mathbb{E}_g fW. \end{aligned}$$

- Hence the importance sampling Algorithm:
- 1. Sample θ_i $i = 1, \ldots, k$ iid from g
  2. Estimate the integral by the unbiased, consistent estimator:

$$\hat{I}_k = \frac{1}{k}\sum_i f(\theta_i)W(\theta_i).$$

- Note that compared to iid Monte Carlo the variance of the estimators changes (typically increases) to $\mathrm{Var}_g(fW)$.

Importance sampling

- Let g be a density such that $\pi(\theta) > 0 \implies g(\theta) > 0$
- Then we can write

$$\begin{aligned} I = \mathbb{E}_\pi f &= \int_\Theta f(\theta)\pi(\theta)d\theta = \int_\Theta f(\theta)\frac{\pi(\theta)}{g(\theta)}g(\theta)d\theta \\ &= \int_\Theta f(\theta)W(\theta)g(\theta)d\theta = \mathbb{E}_g fW. \end{aligned}$$

- Hence the importance sampling Algorithm:
- 1. Sample θ_i $i = 1, \ldots, k$ iid from g
  2. Estimate the integral by the unbiased, consistent estimator:

$$\hat{I}_k = \frac{1}{k}\sum_i f(\theta_i)W(\theta_i).$$

- Note that compared to iid Monte Carlo the variance of the estimators changes (typically increases) to $\mathrm{Var}_g(fW)$.

Importance sampling

- Let g be a density such that $\pi(\theta) > 0 \implies g(\theta) > 0$
- Then we can write

$$
\begin{aligned}
I = \mathbb{E}_\pi f &= \int_\Theta f(\theta)\pi(\theta)d\theta = \int_\Theta f(\theta)\frac{\pi(\theta)}{g(\theta)}g(\theta)d\theta \\
&= \int_\Theta f(\theta)W(\theta)g(\theta)d\theta = \mathbb{E}_g fW.
\end{aligned}
$$

- Hence the importance sampling Algorithm:
- 1. Sample θ_i $i = 1, \ldots, k$ iid from g
  2. Estimate the integral by the unbiased, consistent estimator:

$$\hat{I}_k = \frac{1}{k}\sum_i f(\theta_i)W(\theta_i).$$

- Note that compared to iid Monte Carlo the variance of the estimators changes (typically increases) to $\mathrm{Var}_g(fW)$.

Importance sampling

- Let g be a density such that $\pi(\theta) > 0 \implies g(\theta) > 0$
- Then we can write

$$\begin{aligned} I = \mathbb{E}_\pi f &= \int_\Theta f(\theta)\pi(\theta)d\theta = \int_\Theta f(\theta)\frac{\pi(\theta)}{g(\theta)}g(\theta)d\theta \\ &= \int_\Theta f(\theta)W(\theta)g(\theta)d\theta = \mathbb{E}_g fW. \end{aligned}$$

- Hence the importance sampling Algorithm:
- 1. Sample θ_i $i = 1, \ldots, k$ iid from g
  2. Estimate the integral by the unbiased, consistent estimator:

$$\hat{I}_k = \frac{1}{k}\sum_i f(\theta_i)W(\theta_i).$$

- Note that compared to iid Monte Carlo the variance of the estimators changes (typically increases) to $\mathrm{Var}_g(fW)$.

sequential importance sampling

- The idea can be extended to a Markov process
- if the target distribution is of the form

$$p(\theta_1, \dots, \theta_n) = p(\theta_1) \prod_{i=2}^{n} p(\theta_i|\theta_{i-1})$$

- We can use a proposal process defined by

$$q(\theta_1) \quad \text{and} \quad q(\theta_i|\theta_{i-1}).$$

sequential importance sampling

- The idea can be extended to a Markov process
- if the target distribution is of the form

$$p(\theta_1, \ldots, \theta_n) = p(\theta_1) \prod_{i=2}^{n} p(\theta_i | \theta_{i-1})$$

- We can use a proposal process defined by

$$q(\theta_1) \quad \text{and} \quad q(\theta_i | \theta_{i-1}).$$

sequential importance sampling

- The idea can be extended to a Markov process
- if the target distribution is of the form

$$p(\theta_1, \dots, \theta_n) = p(\theta_1) \prod_{i=2}^{n} p(\theta_i | \theta_{i-1})$$

- We can use a proposal process defined by

$$q(\theta_1) \quad \text{and} \quad q(\theta_i | \theta_{i-1}).$$

sequential importance sampling

- to implement the SIS algorithm:
  1. Sample $\theta_1^{(i)}$ $i = 1, \ldots, k$ iid from q, assign weight
  $$w_1^{(i)} = p(\theta_1^{(i)})/q(\theta_1^{(i)})$$
  2. For $t = 2, \ldots, n$ simulate
  $$\theta_t^{(i)}|\theta_{t-1}^{(i)} \sim q(\theta_t|\theta_{t-1}^{(i)})$$
  and update the weight according to
  $$w_t^{(i)} = w_{t-1}^{(i)} \frac{p(\theta_t^{(i)}|\theta_{t-1}^{(i)})}{q(\theta_t^{(i)}|\theta_{t-1}^{(i)})}$$
- The weakness of importance sampling and SIS is that it is difficult to choose efficient proposal distributions, especially if Θ is high dimensional.

sequential importance sampling

- to implement the SIS algorithm:
  1. Sample $\theta_1^{(i)}$ $i = 1, \ldots, k$ iid from q, assign weight
  $$w_1^{(i)} = p(\theta_1^{(i)})/q(\theta_1^{(i)})$$
  2. For $t = 2, \ldots, n$ simulate
  $$\theta_t^{(i)}|\theta_{t-1}^{(i)} \sim q(\theta_t|\theta_{t-1}^{(i)})$$
  and update the weight according to
  $$w_t^{(i)} = w_{t-1}^{(i)} \frac{p(\theta_t^{(i)}|\theta_{t-1}^{(i)})}{q(\theta_t^{(i)}|\theta_{t-1}^{(i)})}$$
- The weakness of importance sampling and SIS is that it is difficult to choose efficient proposal distributions, especially if Θ is high dimensional.

Markov chains

- Let $P = P(\cdot, \cdot)$ be a Markov operator on a general state space Θ
- This means $P(x, \cdot)$ is a probability measure for every x and for every measurable set A the function $P(\cdot, A)$ is measurable.
- So if
$$\theta_0 \sim \nu$$
then for $t = 1, 2, \ldots$
$$\theta_t \sim P(\theta_{t-1}, \cdot)$$
- The distribution of θ_1 is νP i.e.
$$\nu P(A) = \int_\Theta P(\theta, A)\nu(\theta)d\theta$$
and similarly the distribution of θ_t is νP^t i.e.
$$\nu P^t(A) = \int_\Theta P(\theta, A)\nu P^{t-1}(\theta)d\theta$$

Markov chains

- Let $P = P(\cdot, \cdot)$ be a Markov operator on a general state space Θ
- This means $P(x, \cdot)$ is a probability measure for every x and for every measurable set A the function $P(\cdot, A)$ is measurable.
- So if
$$\theta_0 \sim \nu$$
then for $t = 1, 2, \ldots$
$$\theta_t \sim P(\theta_{t-1}, \cdot)$$
- The distribution of θ_1 is νP i.e.
$$\nu P(A) = \int_\Theta P(\theta, A)\nu(\theta)d\theta$$
and similarly the distribution of θ_t is νP^t i.e.
$$\nu P^t(A) = \int_\Theta P(\theta, A)\nu P^{t-1}(\theta)d\theta$$

Markov chains

- Let $P = P(\cdot, \cdot)$ be a Markov operator on a general state space Θ
- This means $P(x, \cdot)$ is a probability measure for every x and for every measurable set A the function $P(\cdot, A)$ is measurable.
- So if

$$\theta_0 \sim \nu$$

then for $t = 1, 2, \ldots$

$$\theta_t \sim P(\theta_{t-1}, \cdot)$$

- The distribution of θ_1 is νP i.e.

$$\nu P(A) = \int_\Theta P(\theta, A)\nu(\theta)d\theta$$

and similarly the distribution of θ_t is νP^t i.e.

$$\nu P^t(A) = \int_\Theta P(\theta, A)\nu P^{t-1}(\theta)d\theta$$

Markov chains

- Under weak assumptions νP^t converges as $t \to \infty$ to the same measure, say π_{inv} for every initial distribution ν.
- This π_{inv} is called stationary or invariant measure and satisfies for every t
$$\pi_{inv} P^t = \pi_{inv}$$
- So if t is large enough
$$\mathcal{L}(\theta_t) \approx \pi_{inv}$$
- STRATEGY: Take the posterior distribution π and try to design P so that
$$\pi P = \pi.$$
- This is feasible more often than you would expect!!!
- Under very mild conditions this implies
$$\nu P^t \to \pi \quad \text{for every} \quad \nu.$$
- We then have for t large enough approximately
$$\theta_t \sim \pi.$$

Markov chains

- Under weak assumptions νP^t converges as $t \to \infty$ to the same measure, say π_{inv} for every initial distribution ν.
- This π_{inv} is called stationary or invariant measure and satisfies for every t

$$\pi_{inv} P^t = \pi_{inv}$$

- So if t is large enough

$$\mathcal{L}(\theta_t) \approx \pi_{inv}$$

- STRATEGY: Take the posterior distribution π and try to design P so that

$$\pi P = \pi.$$

- This is feasible more often than you would expect!!!
- Under very mild conditions this implies

$$\nu P^t \to \pi \quad \text{for every} \quad \nu.$$

- We then have for t large enough approximately

$$\theta_t \sim \pi.$$

Markov chains

- Under weak assumptions νP^t converges as $t \to \infty$ to the same measure, say π_{inv} for every initial distribution ν.
- This π_{inv} is called stationary or invariant measure and satisfies for every t

$$\pi_{inv} P^t = \pi_{inv}$$

- So if t is large enough

$$\mathcal{L}(\theta_t) \approx \pi_{inv}$$

- STRATEGY: Take the posterior distribution π and try to design P so that

$$\pi P = \pi.$$

- This is feasible more often than you would expect!!!
- Under very mild conditions this implies

$$\nu P^t \to \pi \quad \text{for every} \quad \nu.$$

- We then have for t large enough approximately

$$\theta_t \sim \pi.$$

Markov chains

- Under weak assumptions νP^t converges as $t \to \infty$ to the same measure, say π_{inv} for every initial distribution ν.
- This π_{inv} is called stationary or invariant measure and satisfies for every t

$$\pi_{inv} P^t = \pi_{inv}$$

- So if t is large enough

$$\mathcal{L}(\theta_t) \approx \pi_{inv}$$

- STRATEGY: Take the posterior distribution π and try to design P so that

$$\pi P = \pi.$$

- This is feasible more often than you would expect!!!
- Under very mild conditions this implies

$$\nu P^t \to \pi \quad \text{for every} \quad \nu.$$

- We then have for t large enough approximately

$$\theta_t \sim \pi.$$

Markov chains

- Under weak assumptions νP^t converges as $t \to \infty$ to the same measure, say π_{inv} for every initial distribution ν.
- This π_{inv} is called stationary or invariant measure and satisfies for every t
$$\pi_{inv} P^t = \pi_{inv}$$
- So if t is large enough
$$\mathcal{L}(\theta_t) \approx \pi_{inv}$$
- STRATEGY: Take the posterior distribution π and try to design P so that
$$\pi P = \pi.$$
- This is feasible more often than you would expect!!!
- Under very mild conditions this implies
$$\nu P^t \to \pi \quad \text{for every} \quad \nu.$$
- We then have for t large enough approximately
$$\theta_t \sim \pi.$$

Markov chains

- Under weak assumptions νP^t converges as $t \to \infty$ to the same measure, say π_{inv} for every initial distribution ν.
- This π_{inv} is called stationary or invariant measure and satisfies for every t

$$\pi_{inv} P^t = \pi_{inv}$$

- So if t is large enough

$$\mathcal{L}(\theta_t) \approx \pi_{inv}$$

- STRATEGY: Take the posterior distribution π and try to design P so that

$$\pi P = \pi.$$

- This is feasible more often than you would expect!!!
- Under very mild conditions this implies

$$\nu P^t \to \pi \quad \text{for every} \quad \nu.$$

- We then have for t large enough approximately

$$\theta_t \sim \pi.$$

Markov chains

- Under weak assumptions νP^t converges as $t \to \infty$ to the same measure, say π_{inv} for every initial distribution ν.
- This π_{inv} is called stationary or invariant measure and satisfies for every t
$$\pi_{inv} P^t = \pi_{inv}$$
- So if t is large enough
$$\mathcal{L}(\theta_t) \approx \pi_{inv}$$
- STRATEGY: Take the posterior distribution π and try to design P so that
$$\pi P = \pi.$$
- This is feasible more often than you would expect!!!
- Under very mild conditions this implies
$$\nu P^t \to \pi \quad \text{for every} \quad \nu.$$
- We then have for t large enough approximately
$$\theta_t \sim \pi.$$

CLT for MCMC

- The approach can be validated asymptotically for estimating

$$I(f) = \int_\Theta f(\theta)\pi(\theta)d\theta$$

- if $\theta_0, \theta_1, \ldots$ is a Markov chain with dynamics P, then
- under very mild conditions LLN holds

$$\frac{1}{t}\sum_{i=0}^{t-1} f(\theta_i) \to I(f)$$

- And also under suitable conditions a CLT holds

$$\frac{1}{\sqrt{t}}\sum_{i=0}^{t-1} f(\theta_i) \to N(I(f), \sigma_{as}(P,f))$$

where $\sigma_{as}(P,f)$ is called asymptotic variance.

- There is substantial effort devoted to reliable estimation of $\sigma_{as}(P,f)$.

CLT for MCMC

- The approach can be validated asymptotically for estimating

$$I(f) = \int_\Theta f(\theta)\pi(\theta)d\theta$$

- if $\theta_0, \theta_1, \ldots$ is a Markov chain with dynamics P, then
- under very mild conditions LLN holds

$$\frac{1}{t}\sum_{i=0}^{t-1} f(\theta_i) \to I(f)$$

- And also under suitable conditions a CLT holds

$$\frac{1}{\sqrt{t}}\sum_{i=0}^{t-1} f(\theta_i) \to N(I(f), \sigma_{as}(P,f))$$

where $\sigma_{as}(P,f)$ is called asymptotic variance.
- There is substantial effort devoted to reliable estimation of $\sigma_{as}(P,f)$.

CLT for MCMC

- The approach can be validated asymptotically for estimating

$$I(f) = \int_\Theta f(\theta)\pi(\theta)d\theta$$

- if $\theta_0, \theta_1, \ldots$ is a Markov chain with dynamics P, then
- under very mild conditions LLN holds

$$\frac{1}{t}\sum_{i=0}^{t-1} f(\theta_i) \to I(f)$$

- And also under suitable conditions a CLT holds

$$\frac{1}{\sqrt{t}}\sum_{i=0}^{t-1} f(\theta_i) \to N(I(f), \sigma_{as}(P,f))$$

where $\sigma_{as}(P,f)$ is called asymptotic variance.
- There is substantial effort devoted to reliable estimation of $\sigma_{as}(P,f)$.

CLT for MCMC

- The approach can be validated asymptotically for estimating

$$I(f) = \int_\Theta f(\theta)\pi(\theta)d\theta$$

- if $\theta_0, \theta_1, \ldots$ is a Markov chain with dynamics P, then
- under very mild conditions LLN holds

$$\frac{1}{t}\sum_{i=0}^{t-1} f(\theta_i) \to I(f)$$

- And also under suitable conditions a CLT holds

$$\frac{1}{\sqrt{t}}\sum_{i=0}^{t-1} f(\theta_i) \to N(I(f), \sigma_{as}(P,f))$$

where $\sigma_{as}(P,f)$ is called asymptotic variance.
- There is substantial effort devoted to reliable estimation of $\sigma_{as}(P,f)$.

CLT for MCMC

- The approach can be validated asymptotically for estimating

$$I(f) = \int_{\Theta} f(\theta)\pi(\theta)d\theta$$

- if $\theta_0, \theta_1, \ldots$ is a Markov chain with dynamics P, then
- under very mild conditions LLN holds

$$\frac{1}{t}\sum_{i=0}^{t-1} f(\theta_i) \to I(f)$$

- And also under suitable conditions a CLT holds

$$\frac{1}{\sqrt{t}}\sum_{i=0}^{t-1} f(\theta_i) \to N(I(f), \sigma_{as}(P,f))$$

  where $\sigma_{as}(P,f)$ is called asymptotic variance.
- There is substantial effort devoted to reliable estimation of $\sigma_{as}(P,f)$.

detailed balance and Metropolis Hastings

- One way of ensuring $\pi P = \pi$ is the detailed balance condition

$$\pi(\theta_1)P(\theta_1, \theta_2) = \pi(\theta_2)P(\theta_2, \theta_1)$$

  formally understood as equivalence of measures on $\Theta \times \Theta$.
- In particular consider moving according to some Markov kernel Q
- i.e. from θ_t we propose to move to $\theta_{t+1} \sim Q(\theta_t, \cdot)$
- And this move is accepted with probability $\alpha(\theta_t, \theta_{t+1})$
- Where $\alpha(\theta_t, \theta_{t-1})$ is chosen in such a way that detailed balance holds.
- Many such choices for $\alpha(\theta_t, \theta_{t-1})$ are possible
- One particular (and optimal in a sense beyond the scope of today) is

$$\alpha(\theta_t, \theta_{t+1}) = \min\{1, \frac{\pi(\theta_{t+1})q(\theta_{t+1}, \theta_t)}{\pi(\theta_t)q(\theta_t, \theta_{t+1})}\}.$$

detailed balance and Metropolis Hastings

- One way of ensuring $\pi P = \pi$ is the detailed balance condition

$$\pi(\theta_1)P(\theta_1,\theta_2) = \pi(\theta_2)P(\theta_2,\theta_1)$$

  formally understood as equivalence of measures on $\Theta \times \Theta$.
- In particular consider moving according to some Markov kernel Q
- i.e. from θ_t we propose to move to $\theta_{t+1} \sim Q(\theta_t, \cdot)$
- And this move is accepted with probability $\alpha(\theta_t, \theta_{t+1})$
- Where $\alpha(\theta_t, \theta_{t-1})$ is chosen in such a way that detailed balance holds.
- Many such choices for $\alpha(\theta_t, \theta_{t-1})$ are possible
- One particular (and optimal in a sense beyond the scope of today) is

$$\alpha(\theta_t, \theta_{t+1}) = \min\{1, \frac{\pi(\theta_{t+1})q(\theta_{t+1},\theta_t)}{\pi(\theta_t)q(\theta_t,\theta_{t+1})}\}.$$

detailed balance and Metropolis Hastings

- One way of ensuring $\pi P = \pi$ is the detailed balance condition

$$\pi(\theta_1)P(\theta_1,\theta_2) = \pi(\theta_2)P(\theta_2,\theta_1)$$

  formally understood as equivalence of measures on $\Theta \times \Theta$.
- In particular consider moving according to some Markov kernel Q
- i.e. from θ_t we propose to move to $\theta_{t+1} \sim Q(\theta_t, \cdot)$
- And this move is accepted with probability $\alpha(\theta_t, \theta_{t+1})$
- Where $\alpha(\theta_t, \theta_{t-1})$ is chosen in such a way that detailed balance holds.
- Many such choices for $\alpha(\theta_t, \theta_{t-1})$ are possible
- One particular (and optimal in a sense beyond the scope of today) is

$$\alpha(\theta_t, \theta_{t+1}) = \min\{1, \frac{\pi(\theta_{t+1})q(\theta_{t+1}, \theta_t)}{\pi(\theta_t)q(\theta_t, \theta_{t+1})}\}.$$

detailed balance and Metropolis Hastings

- One way of ensuring $\pi P = \pi$ is the detailed balance condition
$$\pi(\theta_1)P(\theta_1, \theta_2) = \pi(\theta_2)P(\theta_2, \theta_1)$$
formally understood as equivalence of measures on $\Theta \times \Theta$.
- In particular consider moving according to some Markov kernel Q
- i.e. from θ_t we propose to move to $\theta_{t+1} \sim Q(\theta_t, \cdot)$
- And this move is accepted with probability $\alpha(\theta_t, \theta_{t+1})$
- Where $\alpha(\theta_t, \theta_{t-1})$ is chosen in such a way that detailed balance holds.
- Many such choices for $\alpha(\theta_t, \theta_{t-1})$ are possible
- One particular (and optimal in a sense beyond the scope of today) is
$$\alpha(\theta_t, \theta_{t+1}) = \min\{1, \frac{\pi(\theta_{t+1})q(\theta_{t+1}, \theta_t)}{\pi(\theta_t)q(\theta_t, \theta_{t+1})}\}.$$

detailed balance and Metropolis Hastings

- One way of ensuring $\pi P = \pi$ is the detailed balance condition

$$\pi(\theta_1)P(\theta_1, \theta_2) = \pi(\theta_2)P(\theta_2, \theta_1)$$

  formally understood as equivalence of measures on $\Theta \times \Theta$.
- In particular consider moving according to some Markov kernel Q
- i.e. from θ_t we propose to move to $\theta_{t+1} \sim Q(\theta_t, \cdot)$
- And this move is accepted with probability $\alpha(\theta_t, \theta_{t+1})$
- Where $\alpha(\theta_t, \theta_{t-1})$ is chosen in such a way that detailed balance holds.
- Many such choices for $\alpha(\theta_t, \theta_{t-1})$ are possible
- One particular (and optimal in a sense beyond the scope of today) is

$$\alpha(\theta_t, \theta_{t+1}) = \min\{1, \frac{\pi(\theta_{t+1})q(\theta_{t+1}, \theta_t)}{\pi(\theta_t)q(\theta_t, \theta_{t+1})}\}.$$

detailed balance and Metropolis Hastings

- One way of ensuring $\pi P = \pi$ is the detailed balance condition

$$\pi(\theta_1)P(\theta_1,\theta_2) = \pi(\theta_2)P(\theta_2,\theta_1)$$

  formally understood as equivalence of measures on $\Theta \times \Theta$.
- In particular consider moving according to some Markov kernel Q
- i.e. from θ_t we propose to move to $\theta_{t+1} \sim Q(\theta_t,\cdot)$
- And this move is accepted with probability $\alpha(\theta_t,\theta_{t+1})$
- Where $\alpha(\theta_t,\theta_{t-1})$ is chosen in such a way that detailed balance holds.
- Many such choices for $\alpha(\theta_t,\theta_{t-1})$ are possible
- One particular (and optimal in a sense beyond the scope of today) is

$$\alpha(\theta_t,\theta_{t+1}) = \min\{1, \frac{\pi(\theta_{t+1})q(\theta_{t+1},\theta_t)}{\pi(\theta_t)q(\theta_t,\theta_{t+1})}\}.$$

detailed balance and Metropolis Hastings

- One way of ensuring $\pi P = \pi$ is the detailed balance condition

$$\pi(\theta_1)P(\theta_1, \theta_2) = \pi(\theta_2)P(\theta_2, \theta_1)$$

  formally understood as equivalence of measures on $\Theta \times \Theta$.
- In particular consider moving according to some Markov kernel Q
- i.e. from θ_t we propose to move to $\theta_{t+1} \sim Q(\theta_t, \cdot)$
- And this move is accepted with probability $\alpha(\theta_t, \theta_{t+1})$
- Where $\alpha(\theta_t, \theta_{t-1})$ is chosen in such a way that detailed balance holds.
- Many such choices for $\alpha(\theta_t, \theta_{t-1})$ are possible
- One particular (and optimal in a sense beyond the scope of today) is

$$\alpha(\theta_t, \theta_{t+1}) = \min\{1, \frac{\pi(\theta_{t+1})q(\theta_{t+1}, \theta_t)}{\pi(\theta_t)q(\theta_t, \theta_{t+1})}\}.$$

Metropolis-Hastings algorithm

- 1. Given the current state θ_t sample the next step proposal
$$\theta^*_{t+1} \sim Q(\theta_t, \cdot)$$
  2. Set
$$\theta_{t+1} = \theta^*_{t+1} \quad \text{with probability} \quad \alpha(\theta_t, \theta^*_{t+1})$$
  3. Otherwise set $\theta_{t+1} = \theta_t$.
- Exercise: verify the detailed balance for the Metropolis-Hastings algorithm.

The Gibbs Sampler

- For $\Theta = \Theta_1 \times \Theta_2 \times \cdots \times \Theta_d$
- denote the marginals of π as
$$\pi(\theta_k|\theta_{-k})$$
where
$$\theta_{-k} = (\theta_1, \ldots, \theta_{k-1}, \theta_{k+1}, \ldots, \theta_d)$$
- The Gibbs sampler algorithms iterates between updates of
$$\theta_i|\theta_{-i} \sim \pi(\theta_i|\theta_{-i})$$
- There are two basic strategies:
- (1) in each step choosing a coordinate at random (Random Scan Gibbs Sampler)
- (2) Updating systematically one after another (Systematic Scan Gibbs Sampler)
- Literature: Asmussen and Glynn *Stochastic Simulation*

The Gibbs Sampler

- For $\Theta = \Theta_1 \times \Theta_2 \times \cdots \times \Theta_d$
- denote the marginals of π as

$$\pi(\theta_k|\theta_{-k})$$

where

$$\theta_{-k} = (\theta_1, \ldots, \theta_{k-1}, \theta_{k+1}, \ldots, \theta_d)$$

- The Gibbs sampler algorithms iterates between updates of

$$\theta_i|\theta_{-i} \sim \pi(\theta_i|\theta_{-i})$$

- There are two basic strategies:
- (1) in each step choosing a coordinate at random (Random Scan Gibbs Sampler)
- (2) Updating systematically one after another (Systematic Scan Gibbs Sampler)
- Literature: Asmussen and Glynn *Stochastic Simulation*

The Gibbs Sampler

- For $\Theta = \Theta_1 \times \Theta_2 \times \cdots \times \Theta_d$
- denote the marginals of π as
$$\pi(\theta_k|\theta_{-k})$$
where
$$\theta_{-k} = (\theta_1, \ldots, \theta_{k-1}, \theta_{k+1}, \ldots, \theta_d)$$
- The Gibbs sampler algorithms iterates between updates of
$$\theta_i|\theta_{-i} \sim \pi(\theta_i|\theta_{-i})$$
- There are two basic strategies:
- (1) in each step choosing a coordinate at random (Random Scan Gibbs Sampler)
- (2) Updating systematically one after another (Systematic Scan Gibbs Sampler)
- Literature: Asmussen and Glynn *Stochastic Simulation*

The Gibbs Sampler

- For $\Theta = \Theta_1 \times \Theta_2 \times \cdots \times \Theta_d$
- denote the marginals of π as
$$\pi(\theta_k|\theta_{-k})$$
where
$$\theta_{-k} = (\theta_1, \ldots, \theta_{k-1}, \theta_{k+1}, \ldots, \theta_d)$$
- The Gibbs sampler algorithms iterates between updates of
$$\theta_i|\theta_{-i} \sim \pi(\theta_i|\theta_{-i})$$
- There are two basic strategies:
- (1) in each step choosing a coordinate at random (Random Scan Gibbs Sampler)
- (2) Updating systematically one after another (Systematic Scan Gibbs Sampler)
- Literature: Asmussen and Glynn *Stochastic Simulation*

The Gibbs Sampler

- For $\Theta = \Theta_1 \times \Theta_2 \times \cdots \times \Theta_d$
- denote the marginals of π as
$$\pi(\theta_k|\theta_{-k})$$
where
$$\theta_{-k} = (\theta_1, \ldots, \theta_{k-1}, \theta_{k+1}, \ldots, \theta_d)$$
- The Gibbs sampler algorithms iterates between updates of
$$\theta_i|\theta_{-i} \sim \pi(\theta_i|\theta_{-i})$$
- There are two basic strategies:
- (1) in each step choosing a coordinate at random (Random Scan Gibbs Sampler)
- (2) Updating systematically one after another (Systematic Scan Gibbs Sampler)
- Literature: Asmussen and Glynn *Stochastic Simulation*

The Gibbs Sampler

- For $\Theta = \Theta_1 \times \Theta_2 \times \cdots \times \Theta_d$
- denote the marginals of π as
$$\pi(\theta_k|\theta_{-k})$$
where
$$\theta_{-k} = (\theta_1, \ldots, \theta_{k-1}, \theta_{k+1}, \ldots, \theta_d)$$
- The Gibbs sampler algorithms iterates between updates of
$$\theta_i|\theta_{-i} \sim \pi(\theta_i|\theta_{-i})$$
- There are two basic strategies:
- (1) in each step choosing a coordinate at random (Random Scan Gibbs Sampler)
- (2) Updating systematically one after another (Systematic Scan Gibbs Sampler)
- Literature: Asmussen and Glynn *Stochastic Simulation*

The Gibbs Sampler

- For $\Theta = \Theta_1 \times \Theta_2 \times \cdots \times \Theta_d$
- denote the marginals of π as
$$\pi(\theta_k|\theta_{-k})$$
where
$$\theta_{-k} = (\theta_1, \ldots, \theta_{k-1}, \theta_{k+1}, \ldots, \theta_d)$$
- The Gibbs sampler algorithms iterates between updates of
$$\theta_i|\theta_{-i} \sim \pi(\theta_i|\theta_{-i})$$
- There are two basic strategies:
- (1) in each step choosing a coordinate at random (Random Scan Gibbs Sampler)
- (2) Updating systematically one after another (Systematic Scan Gibbs Sampler)
- Literature: Asmussen and Glynn *Stochastic Simulation*